

Divergents_Eval3_Modelling(2)

December 14, 2022

1 Aid Escalating Final Evaluation - 3 (Modelling)

```
[1]: #Loading Important Libraries
```

```
import numpy as np
import pandas as pd
from sklearn.preprocessing import LabelEncoder
import matplotlib.pyplot as plt
import seaborn as sns
import langdetect
import gensim
```

```
[2]: from gensim.models import Doc2Vec
from sklearn import utils
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from gensim.models.doc2vec import TaggedDocument
import re
import seaborn as sns
import matplotlib.pyplot as plt
```

```
[ ]: #Importing necessary libraries
import numpy as np
import pandas as pd
from sklearn.preprocessing import LabelEncoder
from sklearn.feature_extraction.text import TfidfVectorizer
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.utils import shuffle
from sklearn.neighbors import KNeighborsClassifier

import pandas as pd
import numpy as np
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LogisticRegression
```

```

from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score, confusion_matrix
import matplotlib.pyplot as plt
import seaborn as sns

```

```

import nltk
nltk.download('stopwords')
nltk.download('wordnet')
nltk.download('omw-1.4')
import re

```

```

lst_stopwords = nltk.corpus.stopwords.words("english")

```

```

[3]: #Reading Training data
df = pd.read_csv("train.csv")

```

```

[4]: #Reading Test data
df_test = pd.read_csv("test.csv")

```

```

[5]: #Replacing ? with Nan in Train
df.replace("?", np.nan, inplace=True)

```

```

[6]: #Replacing ? with Nan in Test
df_test.replace("?", np.nan, inplace=True)

```

```

[7]: #Total Nan Values
df.isna().sum()

```

```

[7]: link                                0
    link_id                             0
    page_description                     0
    alchemy_category                    1397
    alchemy_category_score              1397
    avg_link_size                       0
    common_word_link_ratio_1            0
    common_word_link_ratio_2            0
    common_word_link_ratio_3            0
    common_word_link_ratio_4            0
    compression_ratio                   0
    embed_ratio                         0
    frame_based                         0
    frame_tag_ratio                     0
    has_domain_link                     0
    html_ratio                          0
    image_ratio                         0

```

is_news	1688
lengthy_link_domain	0
link_word_score	0
news_front_page	727
non_markup_alphanumeric_characters	0
count_of_links	0
number_of_words_in_url	0
parametrized_link_ratio	0
spelling_mistakes_ratio	0
label	0
dtype: int64	

2 NLP Preprocessing

3 Part -1

```
[21]: import re
import nltk

#Appending extra words to Stopwords
lst_stopwords = nltk.corpus.stopwords.words("english")
lst_stopwords.append('title')
lst_stopwords.append('url')
```

```
[22]: # NLP preprocessing to carry out Punctuation Removal, LOvercase Conversion,
↳ Lemmatization, Tokenization

def utils_preprocess_text(text, flg_stemm=False, flg_lemm=True,
↳ lst_stopwords=None):
    ## clean (convert to lowercase and remove punctuations and characters and
↳ then strip)
    text = text.lower()
    # text = re.sub("[\W,\d]", " ", str(text).lower().strip())
    text = re.sub("[\W,\d]", " ", text)

    ## Tokenize (convert from string to list)
    lst_text = text.split()    ## remove Stopwords
    if lst_stopwords is not None:
        lst_text = [word for word in lst_text if word not in
                    lst_stopwords]

    # ## Stemming (remove -ing, -ly, ...)
    # if flg_stemm == True:
    #     ps = nltk.stem.porter.PorterStemmer()
    #     lst_text = [ps.stem(word) for word in lst_text]
```

```

## Lemmatisation (convert the word into root word)
if flg_lemm == True:
    lem = nltk.stem.wordnet.WordNetLemmatizer()

    lst_text = [lem.lemmatize(word) for word in lst_text]
    while " " in lst_text:
        lst_text = lst_text.replace(" ", " ")

    ## back to string from list
    text = " ".join(lst_text)
    return text

```

```

[23]: #Applying previous function on Train data
df["text_clean"] = df["page_description"].apply(lambda x:
↳utils_preprocess_text(x, flg_stemm=False, flg_lemm=True,
↳lst_stopwords=lst_stopwords))

```

```

[24]: ##Applying previous function on Test data
df_test["text_clean"] = df_test["page_description"].apply(lambda x:
↳utils_preprocess_text(x, flg_stemm=False, flg_lemm=True,
↳lst_stopwords=lst_stopwords))

```

```

[25]: # train_text = df["text_clean"]
# test_text = df_test["text_clean"]
# complete_text = pd.concat([df["text_clean"], df_test["text_clean"]])

```

```

[26]: # Train Test Split
train, test = train_test_split(df, test_size=0.3, random_state=0)
import nltk
# train = df1

#Retokenizing text to convert into Doc2vec format
from nltk.corpus import stopwords
def tokenize_text(text):
    tokens = []
    for sent in nltk.sent_tokenize(text):
        for word in nltk.word_tokenize(sent):
            if len(word) < 2:
                continue
            tokens.append(word.lower())
    return tokens
train_tagged = train.apply(lambda r:
↳TaggedDocument(words=tokenize_text(r['text_clean']), tags=[r.label]), axis=1)

```

```
test_tagged = test.apply(lambda r:
    ↳TaggedDocument(words=tokenize_text(r['text_clean']), tags=[r.label]), axis=1)
```

4 Part-2

```
[ ]: # # Remove all punctuations from the text
# import string as st
# def remove_punct(text):
#     return "".join([ch for ch in text if ch not in st.punctuation])
# df["page_description"] = df["page_description"].apply(lambda x:
    ↳remove_punct(x))
# df_test["page_description"] = df_test["page_description"].apply(lambda x:
    ↳remove_punct(x))

# #To remove all digits

# import re
# def rem_digits(text):
#     text = re.sub("\d", " ",text)
#     return text

# df["page_description"] = df["page_description"].apply(lambda x:
    ↳rem_digits(x))
# df_test["page_description"] = df_test["page_description"].apply(lambda x:
    ↳rem_digits(x))

# #Remove url/title related words
# def rem_url(text):
#     text = re.sub("http/https/url/title/www", " ",text)
#     return text
# df["page_description"] = df["page_description"] .apply(lambda x: rem_url(x))
# df_test["page_description"] = df_test["page_description"] .apply(lambda x:
    ↳rem_url(x))

# ''' Convert text to lower case tokens. Here, split() is applied on
    ↳white-spaces. But, it could be applied
#     on special characters, tabs or any other string based on which text is to
    ↳be seperated into tokens.
# '''
# def tokenize(text):
#     text = re.split('\s+',text)
#     return [x.lower() for x in text]
```

```

# import re
# df["page_description"] = df["page_description"].apply(lambda msg :
↳tokenize(msg))
# df_test["page_description"] = df_test["page_description"].apply(lambda msg :
↳tokenize(msg))

# # Remove tokens of length less than 3
# def remove_small_words(text):
#     return [x for x in text if len(x) > 3 and len(x) < 11]

# df["page_description"] = df["page_description"].apply(lambda x :
↳remove_small_words(x))
# df_test["page_description"] = df_test["page_description"].apply(lambda x :
↳remove_small_words(x))

# lst_stopwords = nltk.corpus.stopwords.words("english")
# lst_stopwords.extend(["url", "youre",
↳"dont", "havent", "hadnt", "wont", "wouldnt", "cant", "cannot", "can not", "im",
#
↳"m", "am", "ill", "i will", "its", "it is", "s",
↳"is", "thats", "werent", "doesnt", "didnt", "hasnt"
#
↳"do
↳not", "doesnt", "doesnot", "didnt", "didnot", "hasnt", "hasnot", "havent", "havenot", "hadnt", "wont"
#
↳"wouldnt", "im", "iam", "want", "onto", "into", "www",
↳"url", "http", "https"])
# def remove_stopwords(text):
#     return [word for word in text if word not in lst_stopwords]

# df["page_description"] = df["page_description"].apply(lambda x :
↳remove_stopwords(x))
# df_test["page_description"] = df_test["page_description"].apply(lambda x :
↳remove_stopwords(x))

# # Apply lemmatization on tokens
# def lemmatize(text):
#     word_net = WordNetLemmatizer()
#     return [word_net.lemmatize(word) for word in text]

# df["page_description"] = df["page_description"].apply(lambda x : lemmatize(x))
# df_test["page_description"] = df_test["page_description"].apply(lambda x :
↳lemmatize(x))

# #con

```

```

# m = [df["page_description"][i] for i in range(0, len(df["page_description"]))]
# m1 = [df_test["page_description"][i] for i in range(0,
↳ len(df_test["page_description"]))]
# x = [" ".join(i) for i in m]
# x1 = [" ".join(j) for j in m1]
# x

# df["page_description"] = [i.split(" ")for i in x]
# df_test["page_description"] = [i.split(" ")for i in x1]

# #Unique word Extraction
# def unique(sequence):
#     seen = set()
#     return [x for x in sequence if not (x in seen or seen.add(x))]

# df["page_description"] = df["page_description"].apply(lambda x : unique(x))
# df_test["page_description"] = df_test["page_description"].apply(lambda x :
↳ unique(x))

```

5 Building Vocabulary for Vectorization using Doc2vec

```

[27]: import multiprocessing
cores = multiprocessing.cpu_count()

#Building vocab for Train data(After splitting actual train data)
model_dbow = Doc2Vec(dm=0, vector_size=6000, negative=5, hs=0, min_count=2,
↳ sample = 0, workers=cores, alpha=0.025, min_alpha=0.001)
model_dbow.build_vocab([x for x in tqdm(train_tagged)])

```

100%| | 3105/3105 [00:00<00:00, 3885236.85it/s]

```

[28]: import multiprocessing
cores = multiprocessing.cpu_count()

#Building vocab for Test data with labels (After splitting actual train data)
# model_dbow2 = Doc2Vec(dm=0, vector_size=6000, negative=5, hs=0, min_count=2,
↳ sample = 0, workers=cores, alpha=0.025, min_alpha=0.001)
model_dbow.build_vocab([x for x in tqdm(test_tagged)])

```

100%| | 1332/1332 [00:00<00:00, 2792010.46it/s]

```

[29]: #Tagging each document with its unique key

```

```
final_tagged = df_test.apply(lambda r:
    ↳ TaggedDocument(words=tokenize_text(r['text_clean']), tags = [r.
    ↳ alchemy_category]), axis=1)
```

```
[30]: len(df_test)
```

```
[30]: 2958
```

```
[31]: import multiprocessing
cores = multiprocessing.cpu_count()

#Building vocab for Final tagged data
model_dbow1 = Doc2Vec(dm=0, vector_size=6000, negative=5, hs=0, min_count=2,
    ↳ sample = 0, workers=cores, alpha=0.025, min_alpha=0.001)
model_dbow1.build_vocab([x for x in tqdm(final_tagged)])
```

```
100%|          | 2958/2958 [00:00<00:00, 3600334.08it/s]
```

6 Embedding & Vectorization using Doc2Vec

```
[32]: train_documents = utils.shuffle(train_tagged)

#Function for Vectorization using Doc2vec
model_dbow.train(train_documents, total_examples=len(train_documents), epochs=60)
def vector_for_learning(model, input_docs):
    sents = input_docs
    targets, feature_vectors = zip(*[(doc.tags[0], model.infer_vector(doc.
    ↳ words)) for doc in sents])
    return targets, feature_vectors
```

```
[33]: test_documents = utils.shuffle(test_tagged)
```

```
[34]: # test_documents = utils.shuffle(test_tagged)
# model_dbow2.train(test_documents, total_examples=len(test_documents),
    ↳ epochs=30)
# def vector_for_learning2(model, input_docs):
#     sents = input_docs
#     targets, feature_vectors = zip(*[(doc.tags[0], model.infer_vector(doc.
    ↳ words)) for doc in sents])
#     return targets, feature_vectors
```

```
[35]: test_documents_final= final_tagged
# model_dbow1.
    ↳ train(test_documents_final, total_examples=len(test_documents_final),
    ↳ epochs=30)
# def vector_for_learning1(model, input_docs):
```



```
#     sents = input_docs
#     targets ,feature_vectors = zip(*[(doc.tags[0], model.infer_vector(doc.
    ↪words)) for doc in sents])
#     return targets , feature_vectors
```

```
[36]: # from sklearn.model_selection import StratifiedKFold
# #kf = StratifiedKFold(n_splits=10, shuffle=False)
# search_space = {"penalty":["l1", 'l2', 'elasticnet', None],
#                 "solver": ['liblinear', 'newton-cg', 'newton-cholesky',
    ↪'sag', 'saga'],
#                 "max_iter" : [50,80,100,120],
#
#                 "multi_class" : ['auto', 'ovr', 'multinomial']}]
```

```
[37]: #Main Train data
y_train, X_train = vector_for_learning(model_dbow, train_documents)
```

```
[38]: #Main Test data
y_test, X_test = vector_for_learning(model_dbow, test_documents)
```

```
[40]: #Given test (Unknown label)
y_test1,X_test1 = vector_for_learning(model_dbow,test_documents_final)
```

```
[41]: # importing utility modules
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.metrics import log_loss

# importing machine learning models for prediction

from sklearn.ensemble import RandomForestClassifier
from xgboost import XGBClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.naive_bayes import MultinomialNB

# importing voting classifier
from sklearn.ensemble import VotingClassifier
from sklearn.ensemble import AdaBoostClassifier
from sklearn.model_selection import GridSearchCV
from sklearn.linear_model import RidgeClassifier
```

7 Embedding & Vectorization using TFIDF

```
[59]: # #For Training data

# word_vectorizer = TfidfVectorizer(
#     sublinear_tf = True,
#     strip_accents = 'unicode',
#     analyzer = 'word',
#     token_pattern = '(?u)\b\w+\b\w{,1}',
#     lowercase = False,
#     stop_words = 'english',

#     ngram_range = (1, 1),
#     min_df = 5,
#     max_df = 0.95,
#     norm = 'l2',
#     max_features = 8600
# )
# x = word_vectorizer.fit_transform(m).toarray()
# df_ = pd.DataFrame(x)
# df_

# #For Testing data
# word_vectorizer = TfidfVectorizer(
#     sublinear_tf = True,
#     strip_accents = 'unicode',
#     analyzer = 'word',
#     token_pattern = '(?u)\b\w+\b\w{,1}',
#     lowercase = False,
#     stop_words = 'english',

#     ngram_range = (1, 1),
#     min_df = 5,
#     max_df = 0.95,
#     norm = 'l2',
#     max_features = 8600
# )
# x1 = word_vectorizer.fit_transform(m1).toarray()
# df1_ = pd.DataFrame(x1)
# df1_

# #Concatenating with Label

# # df__ = pd.concat([df_, df["label"]], axis=1)
```

8 Embedding & Vectorization using Word2Vec

```
[60]: # w2v_model = gensim.models.Word2Vec(X_train,
#                                         vector_size=5000,
#                                         window=5,
#                                         min_count=2)

# words = set(w2v_model.wv.index_to_key )
# X_train_vect = np.array([np.array([w2v_model.wv[i] for i in ls if i in words])
#                             for ls in X_train])
# X_test_vect = np.array([np.array([w2v_model.wv[i] for i in ls if i in words])
#                             for ls in X_test])
```

9 1. Modelling using Voting Classifier

```
[42]: # initializing all the model objects with default parameters
model_1 = LogisticRegression(max_iter=2000,
                             multi_class= 'multinomial',
                             penalty= 'l2',
                             solver= 'sag')
model_2 = XGBClassifier(booster='gbtree', callbacks=None,
                       colsample_bylevel=1, colsample_bynode=1, colsample_bytree=0.7,
                       early_stopping_rounds=None, enable_categorical=False,
                       eval_metric=None, feature_types=None, gamma=0.4, eta=0.2,
                       ↪gpu_id=-1, alpha=7,
                       grow_policy='lossguide', importance_type=None,
                       interaction_constraints='', learning_rate=0.05, max_bin=256,
                       max_cat_threshold=64, max_cat_to_onehot=4, max_delta_step=0,
                       max_depth=15, max_leaves=0, min_child_weight=3, missing=None,
                       monotone_constraints='()', n_estimators=300, n_jobs=0,
                       num_parallel_tree=1, predictor='auto', random_state=0,
                       sampling_method='uniform',
                       tree_method='hist', subsample=0.5, )
model_3 = RandomForestClassifier(criterion="entropy", max_depth=500,
                               ↪min_samples_split=10, n_estimators=5000)
#model_4 = AdaBoostClassifier(n_estimators=150, learning_rate=0.
                               ↪01, algorithm="SAMME")
#model_5 = RidgeClassifier(alpha=5.0, fit_intercept=True, copy_X=True,
                               ↪max_iter=5000, tol=0.001, class_weight="balanced", solver='auto')
```

```
[43]: # # Making the final model using voting classifier
# final_model = VotingClassifier(
#     estimators=[('lr', model_1), ('xgb', model_2), ('rf', model_3)],
#     ↪voting='soft', weights = [1.35, 1.67, 0.65])
```

```
[44]: #final_model = VotingClassifier(estimators=[('lr', model_1), ('xgb', model_2),
↳ ('rf', model_3), ('ad', model_4), ('rc', model_5)], voting='hard', weights =
↳ [1.9, 1.7, 1.67, 1.33, 1.45])
```

```
final_model = VotingClassifier(estimators=[('lr', model_1), ('xgb', model_2),
↳ ('rf', model_3)], voting='hard', weights = [1, 1, 2])
```

```
[45]: #Model Fitting
final_model.fit(X_train, y_train)
```

```
[45]: VotingClassifier(estimators=[('lr',
                                   LogisticRegression(max_iter=2000,
                                                         multi_class='multinomial',
                                                         solver='sag')),
                                   ('xgb',
                                   XGBClassifier(alpha=7, base_score=None,
                                                  booster='gbtree', callbacks=None,
                                                  colsample_bylevel=1,
                                                  colsample_bynode=1,
                                                  colsample_bytree=0.7,
                                                  early_stopping_rounds=None,
                                                  enable_categorical=False, eta=0.2,
                                                  eval_metric=None,
                                                  feature_types=None, gamma=0...
                                                  interaction_constraints='',
                                                  learning_rate=0.05, max_bin=256,
                                                  max_cat_threshold=64,
                                                  max_cat_to_onehot=4,
                                                  max_delta_step=0, max_depth=15,
                                                  max_leaves=0, min_child_weight=3,
                                                  missing=None,
                                                  monotone_constraints='()',
                                                  n_estimators=300, n_jobs=0,
                                                  num_parallel_tree=1, ...)),
                                   ('rf',
                                   RandomForestClassifier(criterion='entropy',
                                                           max_depth=500,
                                                           min_samples_split=10,
                                                           n_estimators=5000))],
                        weights=[1, 1, 2])
```

```
[45]: # params = {'voting':['hard', 'soft'],
#               'weights':[(1,1,1,1), (2,1,1,1),
#                           (1,2,1,1), (1,1,2,1),
#                           (1,1,1,2), (1,2,1,1),
#                           (1,1,2,2), (2,1,1,2)]}

# #fit gridsearch & print best params
```

```
# grid = GridSearchCV(final_model, params)
# grid.fit(X_train, y_train)
# print('\n')
# print(f'The best params is : {grid.best_params_}')
```

```
[ ]: # Best params came out to be 'hard' and (1,1,2,1)
```

```
[46]: y_pred_90 = final_model.predict(X_test1)
```

```
[47]: y_pred_90
```

```
[47]: array([1, 0, 0, ..., 0, 0, 0])
```

```
[48]: y_pred_check = final_model.predict(X_test)
```

```
[49]: from sklearn.metrics import accuracy_score, f1_score

print('Testing accuracy %s' % accuracy_score(y_test, y_pred_check))
print('Testing F1 score : {}'.format(f1_score(y_test, y_pred_check,
↪average='weighted')))
```

```
Testing accuracy 0.8040540540540541
Testing F1 score : 0.8027717095518264
```

```
[50]: y_pred_90
```

```
[50]: array([1, 0, 0, ..., 0, 0, 0])
```

```
[51]: x = pd.DataFrame(y_pred_90)
```

```
[52]: x.value_counts()
```

```
[52]: 0    1695
      1    1263
      dtype: int64
```

```
[53]: final_sub = pd.concat([df_test['link_id'],x],axis=1)
```

```
[54]: final_sub.to_csv("final_sub_list.csv")
```

```
[55]: y_pred_t = final_model.predict(X_train)
```

```
[56]: y_pred_t
```

```
[56]: array([1, 1, 1, ..., 1, 0, 0])
```

```
[57]: print('Testing accuracy %s' % accuracy_score(y_train, y_pred_t))
```

Testing accuracy 0.996135265700483

10 2. Modelling using Logistic Regression

```
[ ]: # #Initializing the model
# lr = LogisticRegression()
```

```
[ ]: # from sklearn.model_selection import GridSearchCV
```

```
[ ]: # #Best Params after Hyperparameter tuning
# GS = GridSearchCV(estimator = lr,
#                   param_grid = search_space,
#                   scoring = ["r2", "neg_root_mean_squared_error", "accuracy"],
#                   refit = "r2",
#                   cv = 5,
#                   verbose = 4)
```

```
[ ]: # #Model Fitting
# GS.fit(X_train, y_train)
```

```
[ ]: # #Best Params
# GS.best_params_
```

```
[ ]: # #Best Score
# GS.best_score_
```

```
[29]: # y_train, X_train = vector_for_learning(model_dbow, train_documents)
# y_test, X_test = vector_for_learning(model_dbow, test_documents)

# logreg = LogisticRegression(max_iter=50,
# multi_class= multinomial,
# penalty= 'l1',
# solver= 'saga')
# logreg.fit(X_train, y_train)
# y_pred = logreg.predict(X_test)
# from sklearn.metrics import accuracy_score, f1_score
# print('Testing accuracy %s' % accuracy_score(y_test, y_pred))
# print('Testing F1 score : {}'.format(f1_score(y_test, y_pred,
↪ average='weighted')))
```

Testing accuracy 0.7905405405405406

Testing F1 score : 0.7904502155254269

/home/ibab/.local/lib/python3.8/site-

packages/sklearn/linear_model/_logistic.py:814: ConvergenceWarning: lbfgs failed to converge (status=1):

STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
```

```
[33]: # #Test Prediction
      # y_pred_t = logreg.predict(X_train)
```

```
[34]: # print('Testing accuracy %s' % accuracy_score(y_train, y_pred_t))
```

Testing accuracy 0.879549114331723

```
[35]: # #Vector Learning
      # y_test1,X_test1 = vector_for_learning(model_dbow,test_documents_final)
```

```
[36]: # # Final prediction
      # y_pred_final = logreg.predict(X_test1)
```

```
[37]: # y_pred_final
```

```
[37]: array([1, 0, 0, ..., 0, 0, 0])
```

```
[38]: # x = pd.DataFrame(y_pred_final)
```

```
[40]: # x.value_counts()
```

```
[40]: 0    1569
      1    1389
      dtype: int64
```

```
[42]: # final_sub = pd.concat([df_test['link_id'],x],axis=1)
```

```
[43]: # final_sub
```

```
[43]:   link_id  0
0      4049  1
1      3692  0
2      9739  0
3      1548  1
4      5574  1
...      ...
2953    4257  0
2954   10236  0
2955    5494  0
```

```
2956      9302  0
2957      2633  0
```

```
[2958 rows x 2 columns]
```

```
[44]: # final_sub.to_csv('final_sub_4.csv')
```

11 3. Modelling using XGBClassifier

```
[45]: # from xgboost import XGBClassifier
      # model = XGBClassifier()
      # model.fit(X_train, y_train)
```

```
[45]: XGBClassifier(base_score=0.5, booster='gbtree', callbacks=None,
                  colsample_bylevel=1, colsample_bynode=1, colsample_bytree=1,
                  early_stopping_rounds=None, enable_categorical=False,
                  eval_metric=None, feature_types=None, gamma=0, gpu_id=-1,
                  grow_policy='depthwise', importance_type=None,
                  interaction_constraints='', learning_rate=0.300000012,
                  max_bin=256, max_cat_threshold=64, max_cat_to_onehot=4,
                  max_delta_step=0, max_depth=6, max_leaves=0, min_child_weight=1,
                  missing=nan, monotone_constraints=(), n_estimators=100,
                  n_jobs=0, num_parallel_tree=1, predictor='auto', random_state=0,
                  ...)
```

```
[47]: # y_pred1 = model.predict(X_test)
      # from sklearn.metrics import accuracy_score, f1_score
      # print('Testing accuracy %s' % accuracy_score(y_test, y_pred1))
      # print('Testing F1 score : {}'.format(f1_score(y_test, y_pred1,
      # ↪average='weighted')))
```

```
Testing accuracy 0.7912912912912913
Testing F1 score : 0.7913294116033841
```

```
[50]: # y_pred_final_1 = model.predict(X_test1)
```

```
[52]: # x1 = pd.DataFrame(y_pred_final_1)
```

```
[53]: # x1.value_counts()
```

```
[53]: 0      1509
      1      1449
      dtype: int64
```

```
[58]: # final_sub34 = pd.concat([df_test['link_id'],x1],axis=1)
```



```
[59]: # final_sub34.to_csv('final_sub_99.csv')
```

12 3. Modelling using Random Forest

```
[ ]: # from sklearn.ensemble import RandomForestRegressor  
  
# # create regressor object  
# regressor = RandomForestRegressor()
```

```
[ ]: # from sklearn.model_selection import StratifiedKFold  
# kf = StratifiedKFold(n_splits=10, shuffle=False)  
# search_space = {'n_estimators':[100,200,500],  
#                 'criterion':['gini', 'entropy', 'log_loss'],  
#                 'max_depth':[None,50,100],  
#                 'max_features':['sqrt', 'log2', None],  
  
#                 }
```

```
[ ]: # from sklearn.model_selection import GridSearchCV  
# GS1 = GridSearchCV(estimator = regressor,  
#                    param_grid = search_space,  
#                    scoring = ["r2", "accuracy"],  
#                    refit = "r2",  
#                    cv = 5,  
#                    verbose = 4)
```

```
[ ]: # GS1.fit(X_train, y_train)
```

```
[ ]: # # fit the regressor with x and y data  
# regressor.fit(x_train, y_train)
```

```
[ ]: # y_pred1 = model.predict(X_test)  
# from sklearn.metrics import accuracy_score, f1_score  
# print('Testing accuracy %s' % accuracy_score(y_test, y_pred1))  
# print('Testing F1 score : {}'.format(f1_score(y_test, y_pred1, ↵  
#                                     average='weighted')))
```

```
[ ]: # y_pred_final_1 = model.predict(X_test1)
```

```
[ ]: # x1 = pd.DataFrame(y_pred_final_1)
```

13 4. Modelling using AdaBoost Classifier

```
[ ]: # # Load libraries
# from sklearn.ensemble import AdaBoostClassifier

# # Import Support Vector Classifier
# from sklearn.svm import SVC
# #Import scikit-learn metrics module for accuracy calculation
# from sklearn import metrics
# svc=SVC(probability=True, kernel='linear')

# # Create adaboost classifier object
# abc =AdaBoostClassifier(n_estimators=100, base_estimator=svc, learning_rate=0.
→25)

# model_addf = abc.fit(X_train, y_train)
# #Predict the response for test dataset
# y_pred = model_addf.predict(X_test)

# # Model Accuracy, how often is the classifier correct?
# print("Accuracy:",metrics.accuracy_score(y_test, y_pred))
```

```
[ ]:
```

```
[ ]: # # initializing all the model objects with default parameters
# model_1 = LogisticRegression(max_iter=700,
# multi_class= 'multinomial',
# penalty= 'l1',
# solver= 'sag')
# model_2 = XGBClassifier(base_score=0.5, booster='gbtree', callbacks=None,
# colsample_bylevel=0.5, colsample_bynode=0.5, colsample_bytree=1,
# early_stopping_rounds=None, enable_categorical=False,
# eval_metric=None, feature_types=None, gamma=0.4, eta=0.2,
→gpu_id=-1, alpha=7,
# grow_policy='depthwise', importance_type=None,
# interaction_constraints='', learning_rate=0.05, max_bin=256,
# max_cat_threshold=64, max_cat_to_onehot=4, max_delta_step=0,
# max_depth=15, max_leaves=0, min_child_weight=3, missing=None,
# monotone_constraints='()', n_estimators=300, n_jobs=0,
# num_parallel_tree=1, predictor='auto', random_state=0,
# sampling_method='uniform',
# tree_method='hist')
# model_3 = RandomForestClassifier(criterion="entropy", max_depth=100,
→min_samples_split=5, n_estimators=1000)
# #model_4 = AdaBoostClassifier(n_estimators=150, learning_rate=0.
→01, algorithm="SAMME")
```

```
[ ]: # final_model = VotingClassifier(estimators=[('lr', model_1), ('xgb', model_2),  
↳ ('rf', model_3)], voting='hard', weights = [1, 2, 3])
```

```
[ ]:
```